

1 Jumpstarting Microsoft's New Film Department ¶



2 Overview

This project analyzes film data provided by [IMDB \(https://www.imdb.com/\)](https://www.imdb.com/), [Rotten Tomatoes \(https://www.rottentomatoes.com/\)](https://www.rottentomatoes.com/), & [The Numbers \(https://www.the-numbers.com/\)](https://www.the-numbers.com/). Through the following analysis, I draw connections between responses to a film (box office performance & viewer favorability) and essential film features, such as its total runtime, primary genre and budget size.

3 Business Problem

Microsoft has recently decided to create a new movie studio, but as a company they are completely new to movie making. Thus, they are in need of a data-driven analysis to provide some insight on what kinds of movies will perform well & earn lots if made.

4 Data Preparation

4.1 Importing Libraries & Datasets

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
from glob import glob
from scipy import stats
```

```
In [2]: csv_files = glob("zippedData/*.csv.gz")
```

```
In [3]: csv_files_dict = {}  
for filename in csv_files:  
    filename_cleaned = os.path.basename(filename).replace(".csv.gz", "").re  
    filename_df = pd.read_csv(filename, index_col = 0)  
    csv_files_dict[filename_cleaned] = filename_df
```

```
In [4]: csv_files_dict.keys()
```

```
Out[4]: dict_keys(['imdb_title_crew', 'tmdb_movies', 'imdb_title_akas', 'imdb_tit  
le_ratings', 'imdb_name_basics', 'imdb_title_basics', 'tn_movie_budgets',  
'bom_movie_gross', 'imdb_title_principals'])
```

```
In [5]: title_basics_df = csv_files_dict['imdb_title_basics']  
title_ratings_df = csv_files_dict['imdb_title_ratings']  
movie_budgets_df = csv_files_dict['tn_movie_budgets']
```

```
In [6]: rt_reviews_df = pd.read_csv("zippedData/rt.reviews.tsv.gz", delimiter = '\t'  
rt_info_df = pd.read_csv("zippedData/rt.movie_info.tsv.gz", delimiter = '\t')
```

Now, I examine each of the DataFrames, both to inspect the data types and to search for any abnormalities such as null values.

```
In [7]: movie_budgets_df.shape
```

```
Out[7]: (5782, 5)
```

```
In [8]: movie_budgets_df.isna().sum() / len(movie_budgets_df)
```

```
Out[8]: release_date      0.0  
movie                    0.0  
production_budget      0.0  
domestic_gross          0.0  
worldwide_gross         0.0  
dtype: float64
```

```
In [9]: for col in movie_budgets_df:
        print(f"Currently checking values from col: {col}")
        print(f"Top 5 values:\n{movie_budgets_df[col].value_counts(normalize =
        print("-----")
```

Currently checking values from col: release_date

Top 5 values:

Dec 31, 2014	0.004151
Dec 31, 2015	0.003978
Dec 31, 2010	0.002594
Dec 31, 2008	0.002421
Dec 31, 2012	0.002248

Name: release_date, dtype: float64

Currently checking values from col: movie

Top 5 values:

Home	0.000519
King Kong	0.000519
Halloween	0.000519
Cat People	0.000346
The Fog	0.000346

Name: movie, dtype: float64

Currently checking values from col: production_budget

Top 5 values:

\$20,000,000	0.039952
\$10,000,000	0.036666
\$30,000,000	0.030612
\$15,000,000	0.029920
\$25,000,000	0.029575

Name: production_budget, dtype: float64

Currently checking values from col: domestic_gross

Top 5 values:

\$0	0.094777
\$8,000,000	0.001557
\$2,000,000	0.001211
\$7,000,000	0.001211
\$10,000,000	0.001038

Name: domestic_gross, dtype: float64

Currently checking values from col: worldwide_gross

Top 5 values:

\$0	0.063473
\$8,000,000	0.001557
\$2,000,000	0.001038
\$7,000,000	0.001038
\$4,000,000	0.000692

Name: worldwide_gross, dtype: float64

```
In [10]: title_basics_df.shape
```

```
Out[10]: (146144, 5)
```

```
In [11]: title_basics_df.isna().sum() / len(title_basics_df)
```

```
Out[11]: primary_title      0.000000
         original_title    0.000144
         start_year        0.000000
         runtime_minutes   0.217176
         genres            0.037005
         dtype: float64
```

```
In [12]: for col in title_basics_df:
         print(f"Currently checking values from col: {col}")
         print(f"Top 5 values:\n{title_basics_df[col].value_counts(normalize = T
         print("-----")
```

Currently checking values from col: primary_title

Top 5 values:

```
The Return      0.000137
Broken          0.000137
Homecoming      0.000109
Alone           0.000109
```

Name: primary_title, dtype: float64

Currently checking values from col: original_title

Top 5 values:

```
Home           0.000123
The Return     0.000116
Alone          0.000089
The Gift       0.000089
```

Name: original_title, dtype: float64

Currently checking values from col: start_year

Top 5 values:

```
2016      0.118185
2018      0.115290
2015      0.111144
2014      0.106669
```

Name: start_year, dtype: float64

Currently checking values from col: runtime_minutes

Top 5 values:

```
Series([], Name: runtime_minutes, dtype: float64)
```

Currently checking values from col: genres

Top 5 values:

```
Drama          0.152669
Comedy          0.065207
Horror          0.031065
Comedy,Drama    0.025004
```

Name: genres, dtype: float64

```
In [13]: title_ratings_df.shape
```

```
Out[13]: (73856, 2)
```

```
In [14]: title_ratings_df.isna().sum() / len(title_ratings_df)
```

```
Out[14]: averagerating    0.0
         numvotes         0.0
         dtype: float64
```

```
In [15]: for col in title_ratings_df:
         print(f"Currently checking values from col: {col}")
         print(f"Top 5 values:\n{title_ratings_df[col].value_counts(normalize =
         print("-----")
```

Currently checking values from col: averagerating

Top 5 values:

Series([], Name: averagerating, dtype: float64)

Currently checking values from col: numvotes

Top 5 values:

5 0.036544

7 0.033525

8 0.029341

9 0.026118

Name: numvotes, dtype: float64

```
In [16]: rt_info_df.shape
```

```
Out[16]: (1560, 12)
```

```
In [17]: rt_info_df.isna().sum() / len(rt_info_df)
```

```
Out[17]: id                0.000000
         synopsis          0.039744
         rating            0.001923
         genre             0.005128
         director          0.127564
         writer            0.287821
         theater_date      0.230128
         dvd_date          0.230128
         currency          0.782051
         box_office        0.782051
         runtime           0.019231
         studio            0.683333
         dtype: float64
```

```
In [18]: for col in rt_info_df.drop(columns='synopsis'):
          print(f"Currently checking values from col: {col}")
          print(f"Top 5 values:\n{rt_info_df[col].value_counts(normalize = True)}")
          print("-----")
```

Currently checking values from col: id

Top 5 values:

2000	0.000641
697	0.000641
673	0.000641
674	0.000641
675	0.000641

Name: id, dtype: float64

Currently checking values from col: rating

Top 5 values:

R	0.334618
NR	0.323057
PG	0.154143
PG-13	0.150931
G	0.036609

Name: rating, dtype: float64

Currently checking values from col: genre

Top 5 values:

Drama	0.097294
Comedy	0.070876
Comedy Drama	0.051546
Drama Mystery and Suspense	0.043170
Art House and International Drama	0.039948

Name: genre, dtype: float64

Currently checking values from col: director

Top 5 values:

Steven Spielberg	0.007348
Clint Eastwood	0.005878
William Beaudine	0.002939
Yimou Zhang	0.002939
Bruce Beresford	0.002939

Name: director, dtype: float64

Currently checking values from col: writer

Top 5 values:

Woody Allen	0.0036
Hong Sang-soo	0.0027
John Hughes	0.0027
Jim Jarmusch	0.0027
Sylvester Stallone	0.0027

Name: writer, dtype: float64

Currently checking values from col: theater_date

Top 5 values:

Jan 1, 1987	0.006661
Jan 1, 1994	0.004163
Jan 1, 1973	0.003331
Jun 1, 1990	0.003331
Jan 1, 1966	0.003331

```
Name: theater_date, dtype: float64
-----
Currently checking values from col: dvd_date
Top 5 values:
Jun 1, 2004      0.009159
Sep 3, 2002      0.005828
Nov 6, 2001      0.005828
Aug 27, 1997     0.004996
Sep 2, 2003      0.004996
Name: dvd_date, dtype: float64
-----
Currently checking values from col: currency
Top 5 values:
$      1.0
Name: currency, dtype: float64
-----
Currently checking values from col: box_office
Top 5 values:
32,000,000      0.005882
200,000          0.005882
600,000          0.005882
20,900,803      0.005882
25,957,696      0.002941
Name: box_office, dtype: float64
-----
Currently checking values from col: runtime
Top 5 values:
90 minutes      0.047059
95 minutes      0.043137
100 minutes     0.033333
93 minutes      0.030719
96 minutes      0.028105
Name: runtime, dtype: float64
-----
Currently checking values from col: studio
Top 5 values:
Universal Pictures      0.070850
Paramount Pictures      0.054656
20th Century Fox        0.052632
Sony Pictures Classics   0.044534
Warner Bros. Pictures    0.042510
Name: studio, dtype: float64
-----
```

```
In [19]: rt_reviews_df.shape
```

```
Out[19]: (54432, 8)
```

```
In [20]: rt_reviews_df.isna().sum() / len(rt_reviews_df)
```

```
Out[20]: id          0.000000  
         review      0.102201  
         rating      0.248328  
         fresh       0.000000  
         critic       0.050007  
         top_critic   0.000000  
         publisher    0.005677  
         date         0.000000  
         dtype: float64
```



```
In [21]: for col in rt_reviews_df.drop(columns='review'):
          print(f"Currently checking values from col: {col}")
          print(f"Top 5 values:\n{rt_reviews_df[col].value_counts(normalize = True)}")
          print("-----")
```

Currently checking values from col: id

Top 5 values:

782	0.006210
1067	0.005052
1525	0.004813
1777	0.004777
1083	0.004777

Name: id, dtype: float64

Currently checking values from col: rating

Top 5 values:

3/5	0.105756
4/5	0.089747
3/4	0.087425
2/5	0.077233
2/4	0.066284

Name: rating, dtype: float64

Currently checking values from col: fresh

Top 5 values:

fresh	0.606904
rotten	0.393096

Name: fresh, dtype: float64

Currently checking values from col: critic

Top 5 values:

Emanuel Levy	0.011506
Roger Ebert	0.008915
Dennis Schwartz	0.007987
Nell Minow	0.007194
Frank Swietek	0.006730

Name: critic, dtype: float64

Currently checking values from col: top_critic

Top 5 values:

0	0.759406
1	0.240594

Name: top_critic, dtype: float64

Currently checking values from col: publisher

Top 5 values:

eFilmCritic.com	0.012435
EmanuelLevy.Com	0.010920
New York Times	0.010901
Washington Post	0.010439
Entertainment Weekly	0.009996

Name: publisher, dtype: float64

Currently checking values from col: date

Top 5 values:

January 1, 2000	0.079053
May 20, 2003	0.003711

```
December 6, 2005      0.003325
September 7, 2011    0.002278
July 26, 2002        0.002223
Name: date, dtype: float64
-----
```

4.2 Basic Data Cleaning

Both the **title_basics_df** & **rt_reviews_df** DataFrames have columns (*runtime_minutes* & *reviews*, respectively) containing large chunks of missing values -- about 20% of entries for both.

Additionally, each DataFrame has tens or hundreds of thousands of elements, so both are sufficiently large even after I drop all null values.

```
In [22]: title_basics_df.dropna(inplace=True)
```

```
In [23]: title_basics_df.shape
```

```
Out[23]: (112232, 5)
```

```
In [24]: title_basics_df.isna().sum() / len(title_basics_df)
```

```
Out[24]: primary_title      0.0
         original_title     0.0
         start_year         0.0
         runtime_minutes    0.0
         genres             0.0
         dtype: float64
```

```
In [25]: rt_reviews_df.dropna(inplace=True)
```

```
In [26]: rt_reviews_df.shape
```

```
Out[26]: (33988, 8)
```

```
In [27]: rt_reviews_df.isna().sum() / len(rt_reviews_df)
```

```
Out[27]: id                0.0
         review            0.0
         rating            0.0
         fresh            0.0
         critic            0.0
         top_critic       0.0
         publisher        0.0
         date             0.0
         dtype: float64
```

The columns for *budget* and both *gross* values are not stored as numerical data in the **movie_budgets_df** DataFrame, which will cause issues for analysis.

```
In [28]: movie_budgets_df.head()
```

```
Out[28]:
```

	release_date	movie	production_budget	domestic_gross	worldwide_gross
id					
1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747

I remove the dollar signs and commas from each column before converting their data types to 'integer'.

```
In [29]: def convert_money_int(df, col):
          df[col] = df[col].str.replace("$", "").str.replace(",", "").astype(int)
          return df
```

```
In [30]: cols_to_convert = ['production_budget', 'domestic_gross', 'worldwide_gross']
          for col in cols_to_convert:
              movie_budgets_df = convert_money_int(movie_budgets_df, col)
```

```
In [31]: movie_budgets_df.head()
```

```
Out[31]:
```

	release_date	movie	production_budget	domestic_gross	worldwide_gross
id					
1	Dec 18, 2009	Avatar	425000000	760507625	2776345279
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875
3	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350
4	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747

5 Analysis

5.1 Movie Genre

Using IMDB's data, I explore the link between a movie's primary genre and its worldwide financial performance. This begins with creating new columns for global net gross, release year & return on investment as a percentage.

```
In [32]: movie_budgets_df['net_world'] = movie_budgets_df['worldwide_gross'] - movie
```

```
In [33]: movie_budgets_df['year'] = pd.to_datetime(movie_budgets_df['release_date'])
```

```
In [34]: movie_budgets_df['roi_%'] = (movie_budgets_df['net_world'] / movie_budgets_
```

```
In [35]: movie_budgets_df.head()
```

Out[35]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross	net_wo
id						
1	Dec 18, 2009	Avatar	425000000	760507625	2776345279	23513
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875	6350
3	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350	-2002
4	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963	10724
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747	9997

I merge DataFrames from IMDB & The Numbers into the new **gross_by_genre_df** and remove duplicate entries.

```
In [36]: gross_by_genre_df = pd.merge(movie_budgets_df, title_basics_df, left_on = [
        right_on = ['primary_title', 'start_year'],
        how = 'inner')
gross_by_genre_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1519 entries, 0 to 1518
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   release_date          1519 non-null   object
1   movie                 1519 non-null   object
2   production_budget     1519 non-null   int64
3   domestic_gross        1519 non-null   int64
4   worldwide_gross       1519 non-null   int64
5   net_world             1519 non-null   int64
6   year                 1519 non-null   int64
7   roi_%                1519 non-null   float64
8   primary_title         1519 non-null   object
9   original_title        1519 non-null   object
10  start_year            1519 non-null   int64
11  runtime_minutes       1519 non-null   float64
12  genres                1519 non-null   object
dtypes: float64(2), int64(6), object(5)
memory usage: 166.1+ KB
```

```
In [37]: gross_by_genre_df.isna().sum() / len(gross_by_genre_df)
```

```
Out[37]: release_date      0.0
movie                  0.0
production_budget     0.0
domestic_gross        0.0
worldwide_gross       0.0
net_world             0.0
year                  0.0
roi_%                 0.0
primary_title         0.0
original_title        0.0
start_year            0.0
runtime_minutes       0.0
genres                0.0
dtype: float64
```

```
In [38]: gross_by_genre_df.drop_duplicates(subset=['release_date', 'movie'], inplace
gross_by_genre_df.shape
```

```
Out[38]: (1470, 13)
```

This process appears to have spotted 49 duplicate entries and dropped all of them. Now, I examine the *genres* column to see if it's currently usable for creating a plot.

```
In [39]: gross_by_genre_df.genres
```

```
Out[39]: 0      Action,Adventure,Fantasy
1      Action,Adventure,Sci-Fi
2      Action,Adventure,Sci-Fi
3      Action,Adventure,Sci-Fi
4      Action,Adventure,Fantasy
...
1514      Drama
1515      Horror,Mystery,Thriller
1516      Crime,Drama,Thriller
1517      Drama,Horror,Thriller
1518      Drama
Name: genres, Length: 1470, dtype: object
```

I create a new *genres_split* column, where each element contains a film's given genres separated within a list.

```
In [40]: gross_by_genre_df['genres_split'] = gross_by_genre_df['genres'].map(lambda x: x.split(','))
gross_by_genre_df.head()
```

```
Out[40]:
```

	release_date	movie	production_budget	domestic_gross	worldwide_gross	net_work
0	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875	635063
1	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350	-200237
2	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963	1072413
3	Apr 27, 2018	Avengers: Infinity War	300000000	678815482	2048134200	1748134
4	Nov 17, 2017	Justice League	300000000	229024295	655945209	355945

It is now easy to create specific columns for a film's primary & secondary genres. I am working under the assumption that IMDB has listed the genres for each film in order of relevance.

After that is done, I drop unnecessary columns and check for any ROI outliers that might distort my attempts at visualizations.

```
In [41]: gross_by_genre_df['genre1'] = gross_by_genre_df['genres_split'].map(lambda
```

```
In [42]: y = 'NA'
gross_by_genre_df['genre2'] = gross_by_genre_df['genres_split'].map(lambda
```

```
In [43]: to_drop = ['primary_title', 'original_title', 'start_year',
                    'genres', 'genres_split', 'release_date']
gross_by_genre_df.drop(columns = to_drop, inplace=True)
```

```
In [44]: gross_by_genre_df.sort_values(by='roi_%', ascending=False).head()
```

Out[44]:

	movie	production_budget	domestic_gross	worldwide_gross	net_world	year
1494	The Gallows	100000	22764410	41656474	41556474	2015
1400	The Devil Inside	1000000	53262945	101759490	100759490	2012
1281	Paranormal Activity 2	3000000	84752907	177512032	174512032	2010
1189	Get Out	5000000	176040665	255367951	250367951	2017
1375	Moonlight	1500000	27854931	65245512	63745512	2016


```
In [45]: df1 = gross_by_genre_df[(np.abs(stats.zscore(gross_by_genre_df['roi_%'])) <
```

Now I am ready to visualize the data.

```
In [46]: fig, ax = plt.subplots(figsize=(18,15))

sns.boxplot(x='genre1', y='roi_%', data=df1, ax=ax)
ax.axhline(y=100)
ax.set_title('Film Genre vs. ROI (%)')
ax.set_xlabel('Film Genre')
ax.set_ylabel('Return on Investment (%)')

plt.savefig("roigenrebox.png")
plt.show()
```



In the boxplot above, I added a horizontal line at the y-value of 100% ROI. This marker indicates the point where a film has earned double its production budget.

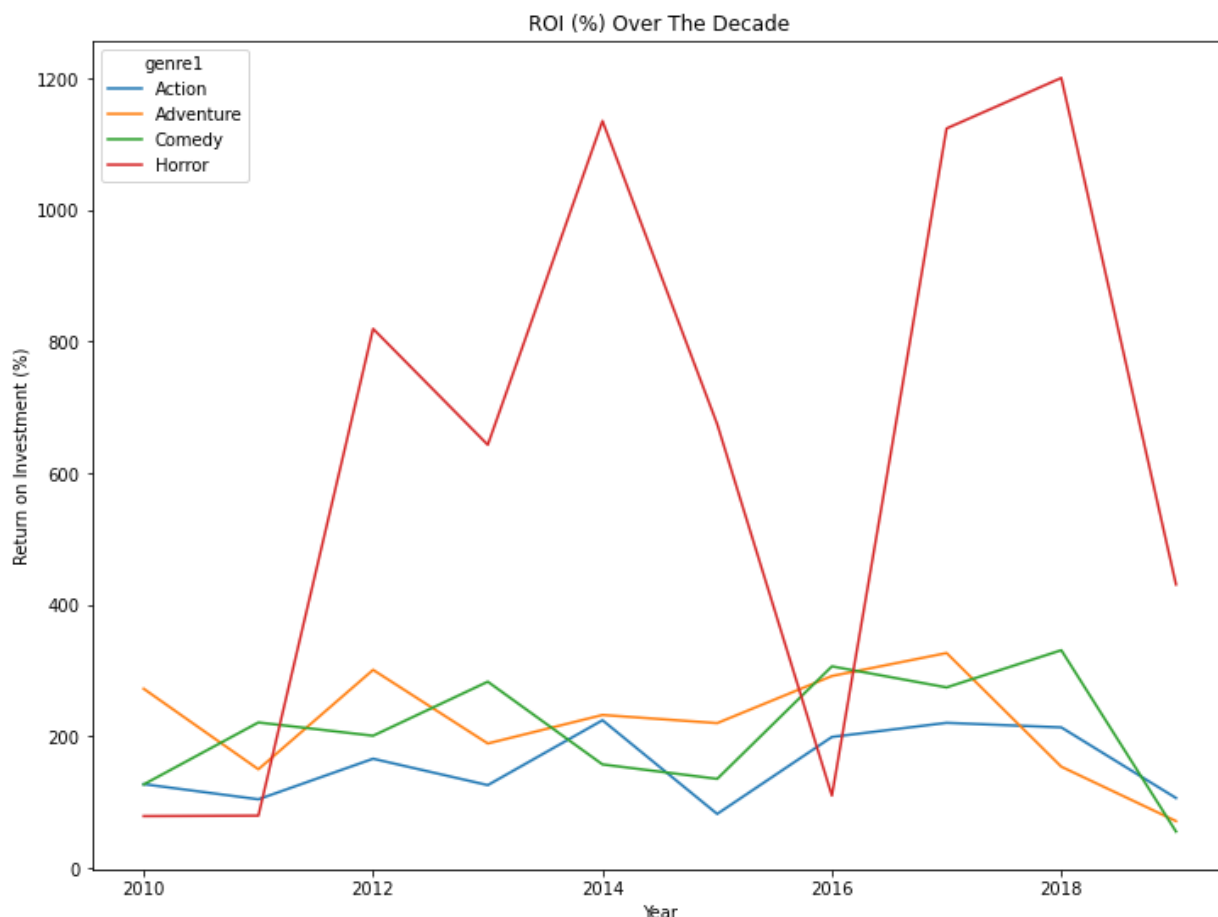
The primary genres of Action, Adventure, Comedy, Horror, Animation & Mystery all have median ROI levels above the 100% threshold. The best looking genres with numerous entries, however, seem to be Action, Adventure, Comedy & Horror, so I will explore the ROI trends for each of these genres further.

```
In [47]: df1 = df1[(df1.genre1 == 'Adventure') | (df1.genre1 == 'Comedy') |  
                (df1.genre1 == 'Horror') | (df1.genre1 == 'Action')]
```

```
In [48]: fig, ax = plt.subplots(figsize=(12,9))

sns.lineplot(x='year', y='roi_%', data=df1, hue='genre1', ci=None, ax=ax)
ax.set_title('ROI (%) Over The Decade')
ax.set_xlabel('Year')
ax.set_ylabel('Return on Investment (%)')

plt.savefig("timeroi.png")
plt.show()
```



It seems as though the advent of streaming has been negatively impacting overall film profitability over the last couple of years, as every genres' ROI percentage has trended downwards since 2018.

Looking at the trends pre-2018, however, reveals that Action, Comedy & Horror movies have made noticeable gains from the beginning of the decade. Action seems to have the most stable, subtle growth, and Comedy has had some noticeable oscillations in profitability. Horror, however, has by far the most erratic jumps up and down in its return on investment.

Adventure has ROI shifts similar to those of Comedy, but overall seems to be trending downwards over the decade, even before 2018.

Based on these figures, I would recommend that Microsoft begin their venture by focusing on primary genres of Action, Comedy & Horror. They should keep in mind, too, that Action might be the most dependable of the three profit-wise.


5.2 Production Budget

The production budgets of films listed by The Numbers have a mean of ~32 million dollars and a median of 17 million dollars. This discrepancy makes it evident that there is significant skew in the budgets.

```
In [49]: movie_budgets_df['production_budget'].describe()
```

```
Out[49]: count      5.782000e+03  
mean       3.158776e+07  
std        4.181208e+07  
min        1.100000e+03  
25%        5.000000e+06  
50%        1.700000e+07  
75%        4.000000e+07  
max        4.250000e+08  
Name: production_budget, dtype: float64
```

```
In [50]: plt.figure(figsize=(6,3))  
sns.histplot(data=movie_budgets_df['production_budget'])  
  
plt.title('Production Budget Count')  
plt.xlabel('Production Budget (Hundreds of Million $)')  
  
plt.savefig("budgetcount.png")
```



It is now especially clear that the collection of budgets is both skewed and enormously wide-ranging. In other words, a single plot that captures every value won't be particularly insightful. Instead, I divide the budgets into different ranges before observing the financial performance of each group.

The most sensible way to categorize production budgets is by using the quantiles shown above. This provides 4 distinct budget tiers -- High Budget, Medium-Low Budget, Medium-High Budget and High Budget.

```
In [51]: bins = [0, 5000000, 17000000, 40000000, np.inf]  
names = ['Low', 'MedLo', 'MedHi', 'High']
```

```
In [52]: movie_budgets_df['budget_tier'] = pd.cut(movie_budgets_df['production_budge
```

```
In [53]: movie_budgets_df.head()
```

Out[53]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross	net_wo
id						
1	Dec 18, 2009	Avatar	425000000	760507625	2776345279	23513
2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875	6350
3	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350	-2002
4	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963	10724
5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747	9997

I inspect the ROI values from **movie_budgets_df** to see, like in the last analysis, if any outliers exist and need to be dropped.

```
In [54]: print(movie_budgets_df.sort_values(by='roi_%', ascending=False).head())
```

	release_date	movie	production_budget	domestic_gr
id				
46	Jun 30, 1972	Deep Throat	25000	45000
14	Mar 21, 1980	Mad Max	200000	8750
93	Sep 25, 2009	Paranormal Activity	450000	107918
80	Jul 10, 2015	The Gallows	100000	22764
7	Jul 14, 1999	The Blair Witch Project	600000	140539

	worldwide_gross	net_world	year	roi_%	budget_tier
id					
46	45000000	44975000	1972	179900.000000	Low
14	99750000	99550000	1980	49775.000000	Low
93	194183034	193733034	2009	43051.785333	Low
80	41656474	41556474	2015	41556.474000	Low
7	248300000	247700000	1999	41283.333333	Low

```
In [55]: print('Mean:', movie_budgets_df['roi_%'].mean())
print('Standard Dev:', movie_budgets_df['roi_%'].std())
print('Upper Outliers Above:', movie_budgets_df['roi_%'].mean() + 3*(movie_
```

```
Mean: 380.01613657949645
Standard Dev: 2953.0282308933056
Upper Outliers Above: 9239.100829259412
```

```
In [56]: movie_budgets_df = movie_budgets_df[movie_budgets_df['roi_%'] <= 9239.1]
print(movie_budgets_df.sort_values(by='roi_%', ascending=False).head())
```

	release_date	movie	production_budget	domestic_gross
id				
76	Feb 15, 1950	Cinderella	2900000	85000000
7	Nov 19, 1925	The Big Parade	245000	11000000
60	Apr 23, 2009	Home	500000	15433
57	Oct 29, 2004	Saw	1200000	55968727
26	Apr 15, 1983	The Evil Dead	375000	2400000

	worldwide_gross	net_world	year	roi_%	budget_tier
id					
76	263591415	260691415	1950	8989.359138	Low
7	22000000	21755000	1925	8879.591837	Low
60	44793168	44293168	2009	8858.633600	Low
57	103880027	102680027	2004	8556.668917	Low
26	29400000	29025000	1983	7740.000000	Low

The median ROI from each group indicates that the high-budget films (at least 40 million dollars) have the highest relative payoff.

```
In [57]: movie_budgets_df.groupby('budget_tier')['roi_%'].median()
```

```
Out[57]: budget_tier
Low      -12.972350
MedLo     65.099995
MedHi     66.240468
High     117.279631
Name: roi_%, dtype: float64
```

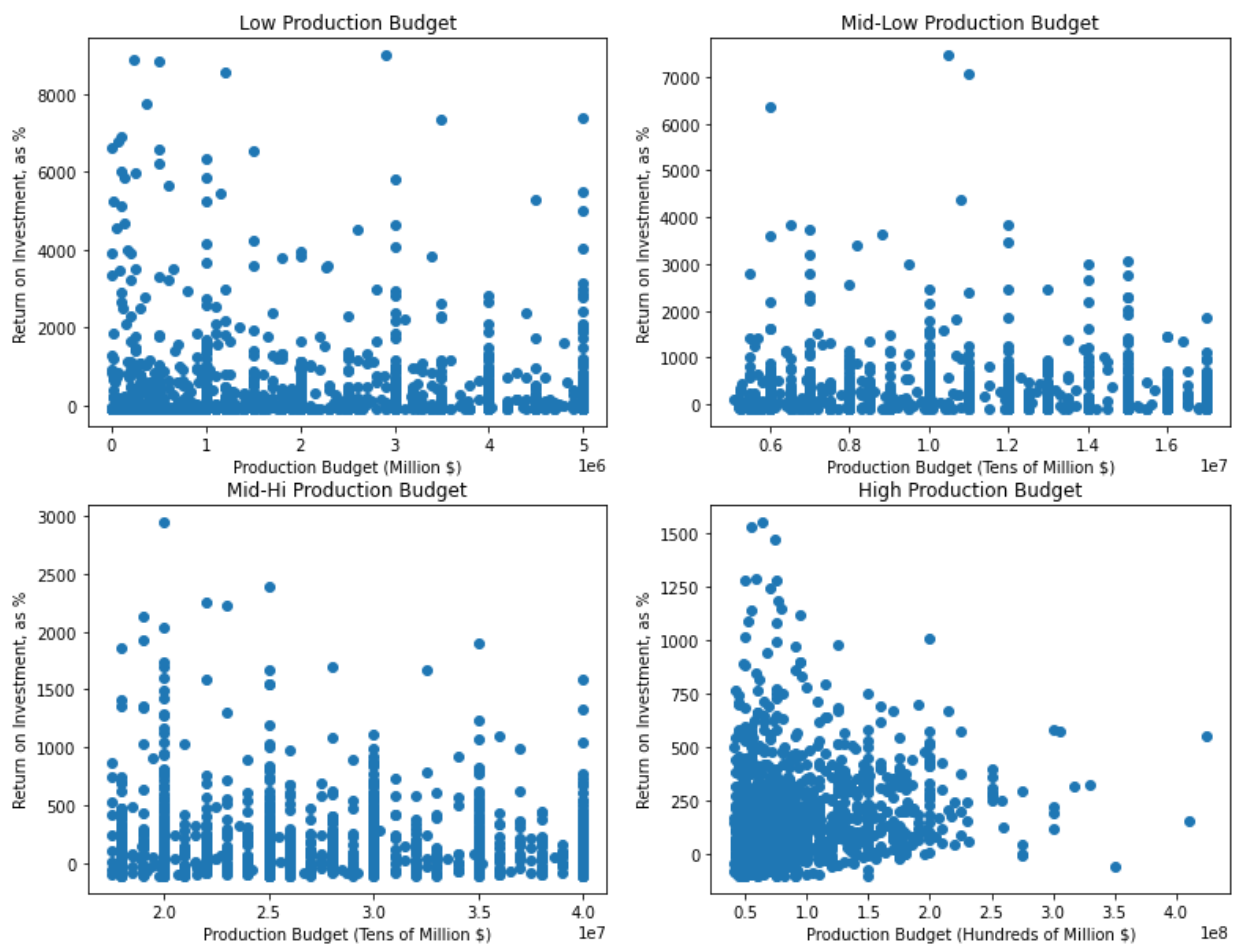
```
In [58]: low_df = movie_budgets_df[movie_budgets_df['budget_tier']=='Low']
medlo_df = movie_budgets_df[movie_budgets_df['budget_tier']=='MedLo']
medhi_df = movie_budgets_df[movie_budgets_df['budget_tier']=='MedHi']
high_df = movie_budgets_df[movie_budgets_df['budget_tier']=='High']
```

```
In [59]: tier_list = [low_df, medlo_df, medhi_df, high_df]
low_df.name = "Low Production Budget"
medlo_df.name = "Mid-Low Production Budget"
medhi_df.name = "Mid-Hi Production Budget"
high_df.name = "High Production Budget"
```

```
In [60]: fig, ax = plt.subplots(nrows=2, ncols=2, figsize=(13,10))

for (i, j) in zip(ax.flatten(), tier_list):
    i.scatter(x= j["production_budget"],
              y= j["roi_%"])
    i.set_title(j.name)
    i.set_ylabel('Return on Investment, as %')
    if j['production_budget'].mean() <= 10000000:
        i.set_xlabel('Production Budget (Million $)')
    elif j['production_budget'].mean() <= 50000000:
        i.set_xlabel('Production Budget (Tens of Million $)')
    else:
        i.set_xlabel('Production Budget (Hundreds of Million $)')

plt.savefig("budgetvsroi.png")
plt.show()
```



(Note that the above plots have different scales for measuring production budget & ROI. Check the axis labels for more detail.)

These plots indicate that there is only a positive correlation between budget & ROI for high budget films. I look at the relationship between budget and global net gross, a more concrete figure, to confirm my conclusion.

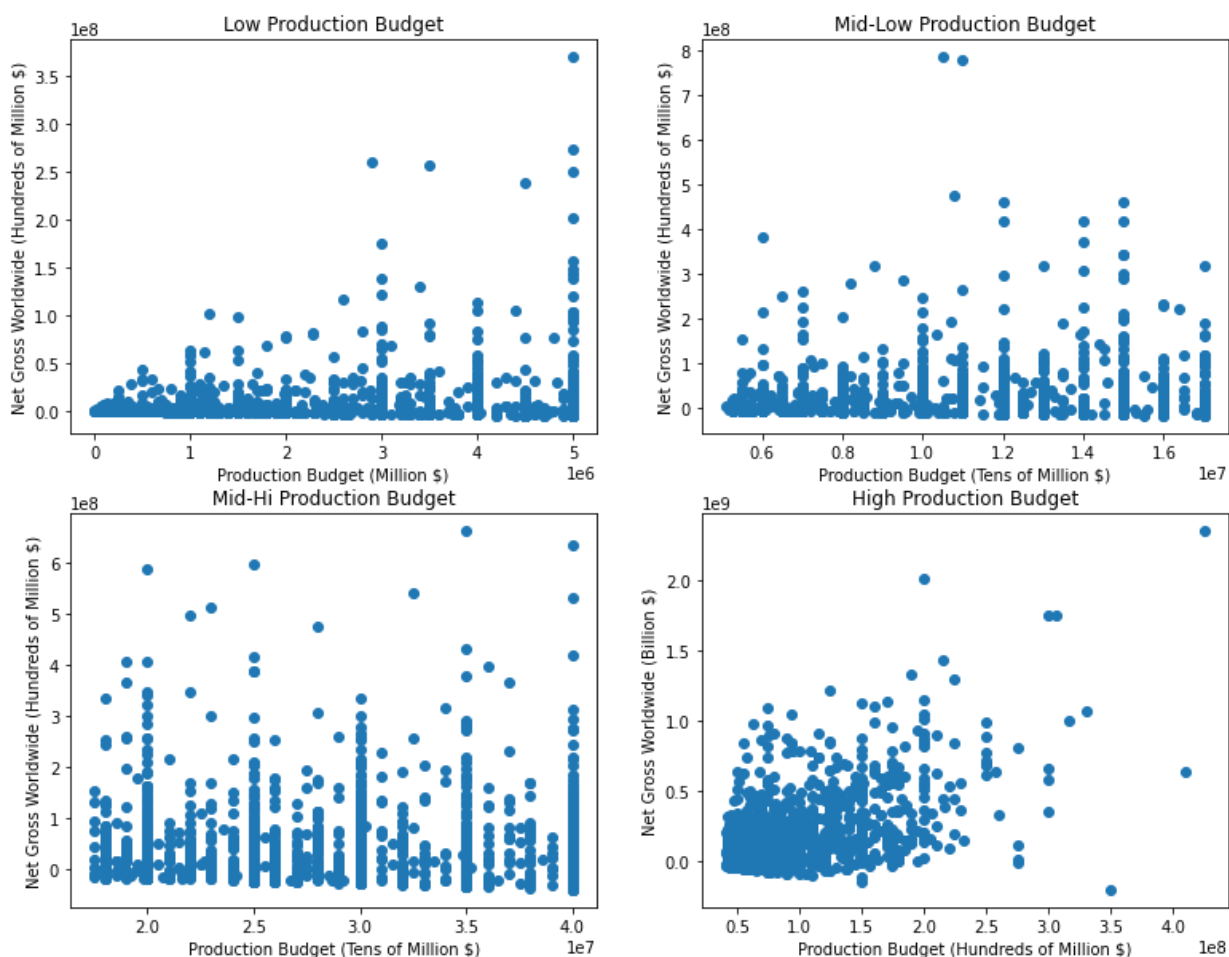
```

In [61]: fig, ax = plt.subplots(nrows=2, ncols=2, figsize=(13,10))

for (i, j) in zip(ax.flatten(), tier_list):
    i.scatter(x= j["production_budget"],
              y= j["net_world"])
    i.set_title(j.name)
    if j['production_budget'].mean() <= 10000000:
        i.set_xlabel('Production Budget (Million $)')
        i.set_ylabel('Net Gross Worldwide (Hundreds of Million $)')
    elif j['production_budget'].mean() <= 50000000:
        i.set_xlabel('Production Budget (Tens of Million $)')
        i.set_ylabel('Net Gross Worldwide (Hundreds of Million $)')
    else:
        i.set_xlabel('Production Budget (Hundreds of Million $)')
        i.set_ylabel('Net Gross Worldwide (Billion $)')

plt.savefig("budgetvsnet.png")
plt.show()

```



These plots further cement my previous interpretation: films with high budgets (40 million USD or more) have by far the most predictable correlation between budget and net financial gain.

5.3 Movie Runtime

I see if a film's runtime affects how positively the audience & critics respond to the film.

5.3.1 First, I inspect ratings_basics_df, which is both IMDB DataFrames merged into one that can display runtime and audience ratings.

```
In [62]: ratings_basics_df = pd.merge(title_basics_df, title_ratings_df, on='tconst')
ratings_basics_df.head()
```

Out[62]:

	primary_title	original_title	start_year	runtime_minutes	genres
tconst					
tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Dram
tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Dram
tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Dram
tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantas
tt0137204	Joe Finds Grace	Joe Finds Grace	2017	83.0	Adventure,Animation,Comed

```
In [63]: ratings_basics_df.drop(columns=['original_title', 'genres'], inplace=True)
ratings_basics_df.drop_duplicates(subset=['primary_title', 'start_year'], i
```

I want to see whether the relationship between runtime and aggregate audience rating is affected at all by the number of votes in a rating. So, as with the previous question, I split the number of votes into four distinct categories by quantile value.

```
In [64]: ratings_basics_df['numvotes'].describe()
```

```
Out[64]: count      6.523600e+04
mean      3.968060e+03
std       3.216720e+04
min       5.000000e+00
25%      1.600000e+01
50%      6.200000e+01
75%     3.530000e+02
max     1.841066e+06
Name: numvotes, dtype: float64
```

```
In [65]: bins2 = [0, 16, 62, 353, np.inf]
names2 = ['LoVotes', 'MedLoVotes', 'MedHiVotes', 'HighVotes']

ratings_basics_df['numvotes_tier'] = pd.cut(ratings_basics_df['numvotes'],
```

```
In [66]: print('DataFrame Dimensions:', ratings_basics_df.shape)
print(ratings_basics_df.head())
```

DataFrame Dimensions: (65236, 6)


	primary_title	start_year	runtime_minutes
\			
tconst			
tt0063540	Sunghursh	2013	175.0
tt0066787	One Day Before the Rainy Season	2019	114.0
tt0069049	The Other Side of the Wind	2018	122.0
tt0100275	The Wandering Soap Opera	2017	80.0
tt0137204	Joe Finds Grace	2017	83.0

	averagerating	numvotes	numvotes_tier
tconst			
tt0063540	7.0	77	MedHiVotes
tt0066787	7.2	43	MedLoVotes
tt0069049	6.9	4517	HighVotes
tt0100275	6.5	119	MedHiVotes
tt0137204	8.1	263	MedHiVotes

I now try making a visual plotting the runtime vs. audience rating, colored by how many votes were cast for the film's rating.

```
In [67]: plt.figure(figsize=(8,6))
sns.scatterplot(data = ratings_basics_df,
                x='runtime_minutes',
                y='averagerating',
                hue='numvotes_tier',
                palette='gist_heat_r'
                )
```

Out[67]: <AxesSubplot:xlabel='runtime_minutes', ylabel='averagerating'>



The plot makes it obvious that there are harmful outliers in the film runtime. I eliminate these through filtering **ratings_basics_df**.

```
In [68]: ratings_basics_df['runtime_minutes'].describe()
```

```
Out[68]: count      65236.000000
mean          94.738595
std          210.141817
min           3.000000
25%          81.000000
50%          91.000000
75%         104.000000
max         51420.000000
Name: runtime_minutes, dtype: float64
```

```
In [69]: runtime = ratings_basics_df['runtime_minutes']
runtime_cleaned_df = ratings_basics_df[runtime < (runtime.mean() + 3*runtime.std())]
print(runtime_cleaned_df.sort_values(by='runtime_minutes', ascending=False))
```

	primary_title	start_year	\
tconst			
tt2261469	Double Fine Adventure	2015	
tt5374716	Chamisso's Shadow	2016	
tt5375100	Paint Drying	2016	
tt8690764	Silence not silence, red not red, live not live	2018	
tt3984388	Close Up	2012	

	runtime_minutes	averagerating	numvotes	numvotes_tier
tconst				
tt2261469	724.0	8.5	59	MedLoVotes
tt5374716	720.0	7.8	19	MedLoVotes
tt5375100	607.0	9.3	218	MedHiVotes
tt8690764	601.0	8.6	22	MedLoVotes
tt3984388	500.0	6.1	13	LoVotes

This helps make the data more plottable. But a look at the new descriptive statistics lets me know that there is still at least one more outlier in need of elimination (see *mean*, *std* & *max* below).

```
In [70]: runtime_cleaned_df['runtime_minutes'].describe()
```

```
Out[70]: count      65230.000000
mean         93.675502
std          24.136521
min           3.000000
25%          81.000000
50%          91.000000
75%         104.000000
max          724.000000
Name: runtime_minutes, dtype: float64
```


```
In [71]: runtime2 = runtime_cleaned_df['runtime_minutes']
runtime_cleaned_df = runtime_cleaned_df[runtime2 < (runtime2.mean() + 3*run
print(runtime_cleaned_df.sort_values(by='runtime_minutes', ascending=False))
```

	primary_title	start_year	runtime_minutes
\			
tconst			
tt7757972	Saakshyam	2018	166.0
tt2309600	Singam 2	2013	166.0
tt2579680	100% Love	2012	166.0
tt2956300	ABCD: American-Born Confused Desi	2013	166.0
tt2320312	Idiot: I Do Ishq Only Tumse	2012	166.0

	averagerating	numvotes	numvotes_tier
tconst			
tt7757972	5.3	262	MedHiVotes
tt2309600	6.3	5046	HighVotes
tt2579680	6.0	296	MedHiVotes
tt2956300	6.7	2141	HighVotes
tt2320312	6.0	10	LoVotes

```
In [72]: plt.figure(figsize=(12,9))
sns.scatterplot(data = runtime_cleaned_df,
                x='runtime_minutes',
                y='averagerating',
                hue='numvotes_tier',
                palette='gist_heat_r'
                )
plt.title('Movie Runtime vs. Audience Rating')
plt.xlabel('Movie Runtime (in minutes)')
plt.ylabel('Aggregate Audience Rating (out of 10)')

plt.savefig("runtimeratingtier.png")
```




The scatter plot is significantly improved, but it is now clear that using all categories of *numvotes_tier* introduces too much noise. I switch to a regplot only using ratings with a "high" number of votes. Of the four categories, it is clearly the closest to representing the actual audience response.

```
In [73]: fig, ax = plt.subplots(figsize=(12,9))

ax.set_title('Movie Runtime vs. Audience Rating (High Vote Count)')
sns.regplot(x='runtime_minutes', y='averagerating',
            data=runtime_cleaned_df[runtime_cleaned_df['numvotes_tier'] ==
            ax=ax,
            color='lightcoral',
            line_kws={'color': 'maroon'},
            fit_reg=True)
ax.set_xlabel('Movie Runtime (in minutes)')
ax.set_ylabel('Aggregate Audience Rating (out of 10)')

plt.savefig("runtimerating.png")
```



This plot suggests a positive relationship between runtime & audience response, especially from the 90 minute runtime onwards. The films with the highest rating floor that still have tight clustering appear around 110-120 minutes.

5.3.2 Next, I inspect review scores from Rotten Tomatoes.

I combine the two Rotten Tomatoes DataFrames into one that give both a film's runtime and its review score. But first, I notice that `rt_reviews_df`'s `rating` column has scores of differing metrics.

```
In [74]: rt_reviews_df.head()
```

Out[74]:

	id	review	rating	fresh	critic	top_critic	publisher	date
0	3	A distinctly gallows take on contemporary fina...	3/5	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018
6	3	Quickly grows repetitive and tiresome, meander...	C	rotten	Eric D. Snider	0	EricDSnider.com	July 17, 2013
7	3	Cronenberg is not a director to be daunted by ...	2/5	rotten	Matt Kelemen	0	Las Vegas CityLife	April 21, 2013
11	3	While not one of Cronenberg's stronger films, ...	B-	fresh	Emanuel Levy	0	EmanuelLevy.Com	February 3, 2013
12	3	Robert Pattinson works mighty hard to make Cos...	2/4	rotten	Christian Toto	0	Big Hollywood	January 15, 2013


```
In [75]: rt_reviews_df['rating'].unique()

Out[75]: array(['3/5', 'C', '2/5', 'B-', '2/4', 'B', '3/4', '4/5', '4/4', '6/10',
                '1/4', '8', '2.5/4', '4/10', '2.0/5', '3/10', '7/10', 'A-', '5/5',
                'F', '3.5/4', 'D+', '1.5/4', '3.5/5', '8/10', 'B+', '9/10',
                '2.5/5', '7.5/10', '5.5/10', 'C-', '1.5/5', '1/5', '5/10', 'C+',
                '0/5', '6', '0.5/4', 'D', '3.1/5', '3/6', '0/4', '2/10', '7', '3',
                'A+', 'A', '4.5/5', '4.0/4', '9.5/10', '2.5', '2.1/2', '6.5/10',
                '3.7/5', '8.4/10', '9', '1', '7.2/10', '2.2/5', '0.5/10', '5',
                '0',
                '2', '1/10', '4.5', '7.7', '5.0/5', '8.5/10', '3.0/5', '0.5/5',
                '1.5/10', '3.0/4', '2.3/10', '4.5/10', '4/6', '3.5', '8.6/10',
                '6/8', 'D-', '2.0/4', '2.7', '4.2/10', '5.8', '4', '7.1/10',
                '3.5/10', '5.8/10', '4.0/5', '0/10', '5.0/10', '5.9/10', '2.4/5',
                '1.9/5', '4.9', '7.4/10', '1.5', '2.3/4', '8.8/10', '4.0/10',
                '2.2', '3.8/10', '6.8/10', '7.3', '7.0/10', '3.2', '4.2', '8.4',
                '5.5/5', '6.3/10', '7.6/10', '8.1/10', '3.6/5', '2/6', '7.7/10',
                '1.8', '8.9/10', '8.9', '8.2/10', '8.3/10', '2.6/6', '4.1/10',
                '2.5/10', 'F+', '6.0/10', '1.0/4', '7.9/10', '8.7/10', '4.3/10',
                '9.6/10', '9.0/10', '4.0', '7.9', '6.7', '8.0/10', '9.2/10', '5.
                2',
                '5.9', '3.7', '4.7', '6.2/10', '1/6', '8.2', '2.6/5', '3.4', '9.
                7',
                '3.3/5', '3.8/5', '1/2', '7.4', '4.8', '1.6/5', '2/2', '1-5',
                '1.0', '4.3/5', '5/6', '9.2', '2.7/5', '4.9/10', '3.0', '3.1',
                '7.8/10', 'F-', '2.3/5', '3.0/10', '3/2', '7.8', '4.2/5', '9.0',
                '7.3/10', '4.4/5', '6.9/10', '0/6', 'T', '6.2', '3.3', '9.8',
                '8.5', '1.0/5', '4.1', '7.1', '3 1/2'], dtype=object)
```

The single number scores and letter grades like 'T' are too vague to be useful and must be dropped. After drops are made, I convert the letter-grade & fraction scores into one standard percent score that doesn't exceed a grade of 100.

```
In [76]: for (index, rating) in rt_reviews_df['rating'].items():
        if '/' not in rating:
            if (rating.isupper() == False):
                rt_reviews_df.drop(index, inplace=True)
            elif rating == 'T':
                rt_reviews_df.drop(index, inplace=True)
            else:
                continue
        elif ' ' in rating:
            rt_reviews_df.drop(index, inplace=True)
        else:
            continue
```

```
In [77]: letter_grade_dict = {'A+': 98, 'A': 95, 'A-': 92, 'B+': 88, 'B': 85, 'B-': 82,
                             'C+': 78, 'C': 75, 'C-': 72, 'D+': 68, 'D': 65, 'D-': 62,
                             'F+': 58, 'F': 55, 'F-': 52}
```

```
In [78]: def score_convert(score):
          if '/' in score:
              fract = score.split('/')
              score = float(fract[0]) * (100.0/float(fract[1]))
          elif score.isupper() == True:
              score = letter_grade_dict[score]
          return score
```

```
In [79]: rt_reviews_df['rating'] = rt_reviews_df['rating'].map(lambda x: score_convert(x))
          rt_reviews_df = rt_reviews_df[rt_reviews_df['rating'] <=100]
```

```
In [80]: rt_reviews_df['rating'].unique()
```

```
Out[80]: array([ 60.          ,  75.          ,  40.          ,  82.          ,
                  50.          ,  85.          ,  80.          , 100.          ,
                  25.          ,  62.5         ,  30.          ,  70.          ,
                  92.          ,  55.          ,  87.5         ,  68.          ,
                  37.5         ,  88.          ,  90.          ,  72.          ,
                  20.          ,  78.          ,   0.          , 12.5         ,
                  65.          ,  62.          ,  98.          ,  95.          ,
                  74.          ,  84.          ,  44.          ,   5.          ,
                  10.          ,  15.          ,  23.          ,  45.          ,
                  66.66666667,  86.          ,  42.          ,  71.          ,
                  35.          ,  58.          ,  59.          ,  48.          ,
                  38.          ,  57.5         ,  63.          ,  76.          ,
                  81.          ,  33.33333333,  77.          ,  89.          ,
                  83.          ,  43.33333333,  41.          ,  79.          ,
                  87.          ,  43.          ,  96.          , 16.66666667,
                  52.          ,  66.          ,  32.          ,  83.33333333,
                  54.          ,  49.          ,  46.          ,  73.          ,
                  69.          ])
```

The two DataFrames **rt_reviews_df** & **rt_info_df** are almost ready to be joined. Before that is done, though, I change the name of the (audience) *rating* column in **rt_reviews_df** to avoid confusion with the (MPA content) *rating* column from **rt_info_df**.

```
In [81]: rt_reviews_df = rt_reviews_df.rename(columns={'rating': 'score'})
```

```
In [82]: info_reviews_df = pd.merge(rt_reviews_df, rt_info_df, on = ['id'], how = 'left')
```

```
In [83]: info_reviews_df.head()
```

```
Out[83]:
```

	id	review	score	fresh	critic	top_critic	publisher	date	synopsi
0	3	A distinctly gallows take on contemporary fina...	60.0	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018	New Y City, r too-dist: future: F
1	3	Quickly grows repetitive and tiresome, meander...	75.0	rotten	Eric D. Snider	0	EricDSnider.com	July 17, 2013	New Y City, r too-dist: future: F
2	3	Cronenberg is not a director to be daunted by ...	40.0	rotten	Matt Kelemen	0	Las Vegas CityLife	April 21, 2013	New Y City, r too-dist: future: F
3	3	While not one of Cronenberg's stronger films, ...	82.0	fresh	Emanuel Levy	0	EmanuelLevy.Com	February 3, 2013	New Y City, r too-dist: future: F
4	3	Robert Pattinson works mighty hard to make Cos...	50.0	rotten	Christian Toto	0	Big Hollywood	January 15, 2013	New Y City, r too-dist: future: F

```
In [84]: cols_to_drop = ['review', 'fresh', 'top_critic', 'date', 'synopsis', 'rating',
                        'genre', 'theater_date', 'dvd_date', 'studio']
info_reviews_df.drop(columns=cols_to_drop, inplace=True)
info_reviews_df.head()
```

Out[84]:

	id	score	critic		publisher	director	writer	currency	box_office
0	3	60.0	PJ Nabarro		Patrick Nabarro	David Cronenberg	David Cronenberg Don DeLillo	\$	600,000
1	3	75.0	Eric D. Snider		EricDSnider.com	David Cronenberg	David Cronenberg Don DeLillo	\$	600,000
2	3	40.0	Matt Kelemen	Las Vegas CityLife		David Cronenberg	David Cronenberg Don DeLillo	\$	600,000
3	3	82.0	Emanuel Levy	EmanuelLevy.Com		David Cronenberg	David Cronenberg Don DeLillo	\$	600,000
4	3	50.0	Christian Toto	Big Hollywood		David Cronenberg	David Cronenberg Don DeLillo	\$	600,000

```
In [85]: print(info_reviews_df.shape)
print(info_reviews_df.isna().sum() / len(info_reviews_df))
```

```
(33334, 9)
id          0.000000
score       0.000000
critic      0.000000
publisher   0.000000
director    0.104698
writer      0.170817
currency    0.303444
box_office  0.303444
runtime     0.017100
dtype: float64
```

```
In [86]: info_reviews_df.runtime.unique()
```

```
Out[86]: array(['108 minutes', '116 minutes', '128 minutes', '95 minutes',
                '82 minutes', '123 minutes', '117 minutes', nan, '90 minutes',
                '97 minutes', '106 minutes', '129 minutes', '98 minutes',
                '127 minutes', '96 minutes', '114 minutes', '110 minutes',
                '75 minutes', '92 minutes', '91 minutes', '100 minutes',
                '99 minutes', '103 minutes', '142 minutes', '119 minutes',
                '122 minutes', '86 minutes', '93 minutes', '111 minutes',
                '89 minutes', '115 minutes', '107 minutes', '165 minutes',
                '113 minutes', '118 minutes', '102 minutes', '101 minutes',
                '135 minutes', '109 minutes', '70 minutes', '124 minutes',
                '105 minutes', '134 minutes', '87 minutes', '188 minutes',
                '104 minutes', '126 minutes', '63 minutes', '132 minutes',
                '137 minutes', '79 minutes', '65 minutes', '147 minutes',
                '171 minutes', '59 minutes', '80 minutes', '94 minutes',
                '88 minutes', '81 minutes', '85 minutes', '130 minutes',
                '143 minutes', '133 minutes', '52 minutes', '83 minutes',
                '121 minutes', '146 minutes', '196 minutes', '78 minutes',
                '120 minutes', '69 minutes', '125 minutes', '153 minutes',
                '145 minutes', '76 minutes', '179 minutes', '64 minutes',
                '138 minutes', '84 minutes', '174 minutes', '144 minutes',
                '184 minutes', '148 minutes', '71 minutes', '358 minutes',
                '67 minutes', '60 minutes', '155 minutes', '150 minutes',
                '164 minutes', '154 minutes', '166 minutes', '140 minutes',
                '112 minutes', '167 minutes', '156 minutes', '170 minutes',
                '205 minutes', '159 minutes', '157 minutes', '136 minutes',
                '58 minutes', '149 minutes', '139 minutes', '161 minutes',
                '56 minutes', '141 minutes', '131 minutes', '229 minutes',
                '180 minutes', '158 minutes', '73 minutes'], dtype=object)
```

I cannot work mathematically or visually with these runtimes as strings, so I remove the "minutes" portion of the string and convert the column to type int. Additionally, I spot a null "nan" label, so any null values from the column must be dropped.

```
In [87]: info_reviews_df.dropna(subset=['runtime'], inplace=True)
info_reviews_df.runtime = info_reviews_df.runtime.str.replace(' minutes', '')
info_reviews_df.dtypes
```

```
Out[87]: id          int64
score         float64
critic         object
publisher      object
director       object
writer         object
currency       object
box_office     object
runtime        int64
dtype: object
```

Because the last collection of runtimes had obvious outliers, I inspect the description of this collection as a precaution to ensure that, at the very least, the max value doesn't sit over 3 standard deviations away from the mean runtime.

```
In [88]: info_reviews_df['runtime'].describe()
```

```
Out[88]: count      32764.00000  
mean         110.43679  
std           20.87213  
min           52.00000  
25%           96.00000  
50%          108.00000  
75%          119.00000  
max           358.00000  
Name: runtime, dtype: float64
```

```
In [89]: runtime3 = info_reviews_df['runtime']  
info_reviews_df = info_reviews_df[runtime3 < (runtime3.mean() + 3*runtime3.  
print(info_reviews_df.describe())
```


	id	score	runtime
count	32529.000000	32529.000000	32529.000000
mean	1047.001014	66.197065	109.581174
std	585.095333	20.914943	17.630648
min	3.000000	0.000000	52.000000
25%	554.000000	50.000000	96.000000
50%	1083.000000	70.000000	108.000000
75%	1545.000000	80.000000	119.000000
max	2000.000000	100.000000	171.000000

Now that the outliers have been filtered out, I am ready to put the plot together.

```
In [90]: fig, ax = plt.subplots(figsize=(12,9))

ax.set_title('Movie Runtime vs. Rotten Tomatoes Scores')
sns.regplot(x='runtime', y='score',
            data=info_reviews_df,
            ax=ax,
            color='turquoise',
            line_kws={'color': 'darkcyan'},
            fit_reg=True)
ax.set_xlabel('Movie Runtime (in minutes)')
ax.set_ylabel('Critic Score (out of 100)')

plt.savefig("runtimescore.png")
```



This dataset allows for each film to have multiple scores associated with it, making the plot a little more tricky to navigate, but nonetheless informative. Overall, there seems to be a slight positive correlation between runtime and critic score.

The higher (60%+) critic scores seem to cluster most tightly between runtimes of 80 & 140 minutes, further solidifying the runtime range previously explored.

6 Conclusions

My analysis leads to the following recommendations for Microsoft as they begin their venture into film:

- **Focus primarily on creating Action, Comedy or Horror films before determining where to branch out.** Not only do these genres have some of the most impressive returns on investment, but each genre's profitability has grown, with varying stability, over most of the last decade. Action movies seem to be the safest bet of the three with regards to profit.
- **Ensure a positive return on investment by making as many high-budget films (at least 40 million USD) as possible.** Only at these large budget levels does there seem to be a positive correlation between the money put into making a film and its relative financial gains. Microsoft is also in the unique position of being able to pay for big budgets from the start.
- **Gain favorability with audiences & critics by releasing movies that run between 100 & 120 minutes.** Looking at the responses provided by IMDB and Rotten Tomatoes, it seems apparent that the films with tightly clustered positive scores and relatively low amounts of negative reactions have runtimes in this range. Additionally, there does appear to be positive correlation between runtime and response, so do not be afraid to test the upper limits of acceptable runtimes.

7 Next Steps

Given the necessary time, additional analysis could provide insight on questions such as:

- Which experienced studios could Microsoft benefit from a partnership with?
- Which screenwriters make the most profitable scripts?
- Which release months or seasons are best for a film's financial performance?

In []: